

## PP-f22

[Assignments](#) / Programming Assignment VI

# Programming Assignment VI: OpenCL Programming

Parallel Programming by Prof. Yi-Ping You

Due date: 23:59, Jan 5, Thursday, 2023

The purpose of this assignment is to familiarize yourself with OpenCL programming.

## TABLE OF CONTENTS

- 1 [Programming Assignment VI: OpenCL Programming](#)
  - a [1. Image convolution using OpenCL](#)
  - b [2. Requirements](#)
  - c [3. Grading Policy](#)
  - d [4. Evaluation Platform](#)
  - e [5. Submission](#)
  - f [6. References](#)

Get the source code:

```
$ wget http://nycu-sslab.github.io/PP-f22/assignments/HW6/HW6.zip
$ unzip HW6.zip -d HW6
$ cd HW6
```

## 1. Image convolution using OpenCL

Convolution is a common operation in image processing. It is used for blurring, sharpening, embossing, edge detection, and more. The image convolution process is accomplished by doing a convolution between a small matrix (which is called a *filter kernel* in image processing) and an image. You may learn more about the convolution process at [Wikipedia: Convolution](#).

Figure 1 shows an illustration of the concept of applying a convolution filter to a specific pixel, value of which is 3. After the convolution process, the value of the pixel becomes 7 —how the resulting value is computed is illustrated on the right of the figure.




Figure 1. Applying a convolution filter to the dark gray pixel of the source image (value of which is 3 ).

In this assignment, you will need to implement a GPU kernel function for convolution in OpenCL by using the zero-padding method. A serial implementation of convolution can be found in `serialConv()` in `serialConv.c`. You can refer to the implementation to port it to OpenCL. You may refer to [this article](#) to learn about the zero-padding method. Figure 2 shows an example of applying the zero-padding method to the source image (on the left) and thereby resulting a same-size, filtered output image (on the right).




Figure 2. Applying the zero-padding method to a source image.

Your job is to parallelize the computation of the convolution using OpenCL. A starter code that spawns OpenCL threads is provided in function `hostFE()`, which is located in `hostFE.c`. `hostFE()` is the host front-end function that allocates memories and launches a GPU kernel, called `convolution()`, which is located in `kernel.cl`.

Currently `hostFE()` and `convolution()` do not do any computation and return immediately. You should complete these two functions to accomplish this assignment.

You can build the program by typing `make`, and run the program via `./conv`. Your program should read an input image from `input.bmp`, perform image convolution, and output the result image into `output.bmp`.

You can use the following command to test your own image:

```
ffmpeg -i source.png -pix_fmt gray -vf scale=600:400 destination.bmp
```

You can use a different filter kernel by adding option `-f N` when running the program (i.e., `./conv -f N`), where `N` is either 1 (by default), 2, or 3, and indicates which filter kernel is used. Each filter kernel is defined in a CSV file (`filter1.csv`, `filter2.csv`, or `filter3.csv`). The first line of the CSV file defines the width (or height) of the filter kernel, and the remaining lines define the values of the filter kernel.

## 2. Requirements

You will modify only `hostFE.c` and `kernel.cl`.

**Q1 (5 points):** Explain your implementation. How do you optimize the performance of convolution?

**[Bonus] Q2 (10 points):** Rewrite the program using CUDA. (1) Explain your CUDA implementation, (2) plot a chart to show the performance difference between using OpenCL and CUDA, and (3) explain the result.

Answer the questions marked with **Q1** (and **Q2**) in a **REPORT** using [HackMD](#). Notice that in this assignment a higher standard will be applied when grading the quality of your report.

**Note:** You cannot print any message in your program.

## 3. Grading Policy

**NO CHEATING!!** You will receive no credit if you are found cheating.

Total of 100+10%:

- Correctness (75%): 25% for each of the three filters ( `filter1.csv` , `filter2.csv` , and `filter3.csv` ). The breakdown of the 25%:
  - 10%: Your parallelized program passes the verification.
  - 15%: The speedup over the serial version should be greater than 5.0 for filter1 and 3, 4.0 for filter2.
- Performance (20%): Compete with your classmates. See the metric below.
- Questions (5+10%).

Metric:

$$\frac{T - Y}{T - F} \times 60\%, \text{ if } Y < T + \begin{cases} 40\%, \text{ if } Y < F \times 2 \\ 20\%, \text{ else} \end{cases}$$

where  $Y$  and  $F$  indicate the execution time of your program and the fastest program, respectively, and  $T = F \times 1.5$ .

## 4. Evaluation Platform

Your program should be able to run on UNIX-like OS platforms. We will evaluate your programs on the workstations dedicated for this course. You can access these workstations by `ssh` with the following information.

The workstations are based on Ubuntu 20.04 with Intel(R) Core(TM) i5-7500 CPU @ 3.40GHz processors and **GeForce GTX 1060 6GB**. `g++-10` , `clang++-11` , `cuda11.5` , and `openCL3.0` have been installed.

IP	Port	User Name	Password
140.113.215.197	10002-10010	{student_id}	{your_passwd}

Login example:

```
$ ssh <student_id>@140.113.215.197 -p <port>
```

You can use the testing script `test_hw6` to check your answer *for reference only*. Run `test_hw6` in a directory that contains your `HW6_XXXXXXX.zip` file on the workstation.

## 5. Submission

All your files should be organized in the following hierarchy and zipped into a `.zip` file, named `HW6_XXXXXXX.zip` , where `XXXXXXX` is your student ID.

Directory structure inside the zipped file:

- HW6\_XXXXXXX.zip (root)
  - kernel.cl
  - hostFE.c
  - url.txt

Notice that you just need to provide the URL of your HackMD report in `url.txt` , and enable the write permission for someone who knows the URL so that TAs can give you feedback directly in your report.

Zip the file:

```
$ zip HW6_XXXXXXX.zip kernel.cl hostFE.c url.txt
```

Be sure to upload your zipped file to new E3 e-Campus system by the due date.

You will get *NO POINT* if your ZIP's name is wrong or the ZIP hierarchy is incorrect.

You will get a 5-point penalty if you hand out unnecessary files (e.g., obj files, .vscode, .\_\_MACOSX).

## 6. References

- [NVIDIA OpenCL SDK Code Samples](#)
- [Getting started with OpenCL and GPU Computing](#)

[Back to top](#)

This website is built using [Kevin Lin's Just the Class](#) Jekyll template.